



Chapter 8

Greedy Algorithms on Graphs

References:

[DPV 4.4-4.5]

[KT 4.4-4.6]

[CLRS 23, 24.3, 6]


[SSS 6.1, 6.3]



Part 1: Minimum Spanning Trees

The Problem

- Given a set of cities, how to connect them by road networks?
 - Result: all cities are reachable one-another
 - Shortest **total** length of roads
 - (does not necessarily give shortest length between a given pair of cities)
- Abstract model
 - Given: a connected undirected graph G
 - Edges have weights (e.g. distance)
 - Goal: a set of edges that connect all vertices with minimum total weight





Minimum Spanning Trees

- A *spanning tree* of an undirected graph is a subset of edges such that the graph remains connected
 - Must be a tree (i.e. no cycles)
 - Reason: if there is a cycle, we can remove one edge in the cycle and all vertices remain connected
- A *minimum spanning tree* (MST) is a spanning tree with the minimum total edge weight
- How to find the MST?
 - Greedy algorithms?



Greedy Algorithm for MST


- Idea: repeatedly add the shortest unused edge
 - As long as it connects two different components, i.e. not redundant
- This is *Kruskal's algorithm*
 - High-level pseudocode: (many details to be filled in later)

```
MST-Kruskal(G) {  
    E := sorted list of edges (by weight)  
    T := G without edges  
    while T has fewer than n-1 edges {  
        Remove first edge e = (u,v) from E  
        if (u, v in different components in T)  
            add e to T  
    }  
}
```

Kruskal's Algorithm: Example


- Edge labels = weights

- (a, g): add
- (c, d): add
- (b, g): add
- (c, g): add
- (a, b): cycle, not add
- (b, c): cycle, not add
- (d, g): cycle, not add
- (d, e): add
- (e, g): cycle, not add
- (a, f): add




Kruskal's Algorithm: Optimality

- Does it always produce the MST?
 - Certainly it is a spanning tree; is it minimum?
- Proof by contradiction. Sketch:
- Assume there is an edge e in the true MST T^* but not in the tree produced by Kruskal's algorithm T^K
- Edge e separates T^* into two parts, $S1$ and $S2$
- Since e is not in T^K , T^K connects x and y by another path of edges



Proof of Optimality (cont'd)

- In this path P , there is at least one edge e' connecting $S1$ and $S2$
- e' must have smaller weight than e since T^K adds edges by increasing order of weights
- Hence we can remove e from T^* and add e' to T^* to get another tree with smaller weight
- This contradicts optimality of T^*






Kruskal's Algorithm: Running Time

- Sorting edge weights: $O(m \log m)$
 - $= O(m \log n)$ (recall $m = O(n^2)$)
- Inside while loop: depends on two important operation
 - 1) Checking whether two vertices are already connected (i.e. in the same component)
 - At most $O(m)$ times in total, over all executions of loop
 - 2) Connecting two vertices (components) together when an edge is added
 - At most $O(n)$ times in total, over all executions of loop (since the final MST has $n - 1$ edges only)
- How to perform these operations efficiently?

Kruskal's using Disjoint Sets

- Representing connected components as *disjoint sets*
 - Initially: $\{a\} \{b\} \{c\} \{d\} \{e\} \{f\} \{g\}$
 - Add (a,g): $\{a,g\} \{b\} \{c\} \{d\} \{e\} \{f\}$
 - Add (c,d): $\{a,g\} \{b\} \{c,d\} \{e\} \{f\}$
 - Add (b,g): $\{a,b,g\}, \{c,d\} \{e\} \{f\}$
 - Add (c,g): $\{a,b,c,d,g\} \{e\} \{f\}$
 - Do not add (a,b) [same set]
 - Do not add (b,c)
 - Do not add (d,g)
 - Add (d,e): $\{a,b,c,d,e,g\} \{f\}$
 - Do not add (e,g)
 - Add (a,f): $\{a,b,c,d,e,f,g\}$





Disjoint Set Data Structure

- We want a data structure for storing:
 - A number of disjoint sets
 - Each set contains elements from the same ground set
- And support the following operations efficiently:
 - **Find(u)**: given an element u , return a “name” of the set containing that element
 - A simple name: some element in the set (“set representatives”)
 - **Union(u, v)**: merge two sets containing u and v
 - Often, u and v are already the set representatives



Using Union-find in MST Algorithms

- Find() and Union() operations support the MST algorithm:
 - Each vertex begins as an individual set initially
 - Edge joining vertices → merge into same component → merge the disjoint sets
 - To check whether two vertices u, v belong to same component: just check whether $\text{Find}(u) = \text{Find}(v)$
 - To merge two vertices u and v , find their set representatives, and then merge: $\text{Union}(\text{Find}(u), \text{Find}(v))$
- We need data structures supporting the union and find operations



Union-find Try #1: Array

- Use an array to keep track of the set names

- $A[i]$ = set name of element i

- Example: $\{1,2,7\}$ $\{3,4\}$ $\{5\}$ $\{6\}$

$A[i]$

1	1	3	3	5	6	1
---	---	---	---	---	---	---

- **Find(i)**: just return $A[i]$

- $O(1)$ time

- **Union(i, j)**: find the set names of i and j , change all array entries with those two names to the same name

- $O(n)$ time

- Example: $\text{union}(2, 5)$


$A[i]$

1	1	3	3	1	6	1
---	---	---	---	---	---	---



Union-find Try #2: Linked List

- Put all objects in same disjoint set (same name) as a list. Head of list is set representative
- **Find(u)**: $O(n)$ time (go through all lists)
- **Union(u,v)**: can be done in $O(1)$ time if use doubly-circular linked list





Union-find #3: Array-and-Linked-List

- For each element, keep track of:
 - Set name
 - Pointer of next element in the set
 - Size of the set
- Size of set is kept because we use the *weighted-union heuristic*: always modify the smaller set
- Example:

Initially: {1} {2} {3} {4} {5} {6} {7}

1	2	3	4	5	6	7
0	0	0	0	0	0	0
1	1	1	1	1	1	1

← Set name

← Next element of the set,
0 = null

← Size of set

Union-find Example (Cont'd)

Union(1,7): {1,7} {2} {3} {4} {5} {6}

1	2	3	4	5	6	1
7	0	0	0	0	0	0
2	1	1	1	1	1	-

1's next element
is 7

7's name is now 1

Size of {1,7} is 2

Don't need to keep size anymore

Union(3,4): {1,7} {2} {3,4} {5} {6}

1	2	3	3	5	6	1
7	0	4	0	0	0	0
2	1	2	-	1	1	-



Union-find Example (Cont'd)

Union(2,7): {1,2,7} {3,4} {5} {6}

1	1	3	3	5	6	1
2	7	4	0	0	0	0
3	-	2	-	1	1	-

{1,7} has size 2 while {2} has size 1, so only change set name of {2}

Union(3,7): {1,2,3,4,7} {5} {6}

1	1	1	1	5	6	1
3	7	4	2	0	0	0
5	-	-	-	1	1	-

Follow this virtual "linked list" to find entries to modify

Union(1,2): no change (Set names of 1,2 the same)

Union(2,3): no change... and so on




Running Time of Kruskal's

- Using the array-and-linked-list data structure:
 - Find(u) in $O(1)$ time
 - Union(u, v) is still worst-case $O(n)$ time (update entries)
- Better analysis for Union():
 - Every time a union occurs, only the smaller set is modified
 - After modification, size of modified set at least doubled
 - Therefore, each entry can be modified at most $\log n$ times
 - Total time for n Union() = $O(n \log n)$
- Overall running time
 - $O(m \log n)$ for sorting, $O(m)$ for Find, $O(n \log n)$ for Union
 - Total $O(m \log n)$
 - There are better data structures with faster running times

Another Approach?

- Is there another way of “greedily” adding edges?
- Idea: “grow” the component from a vertex
- Only one partial MST when the algorithm is running, not a forest (as in Kruskal’s)
- Example: (starting from a)





Prim's Algorithm

- Algorithm:
 - At every step, find minimum-weight outgoing edge
 - Enlarge component
- Finding minimum outgoing edge:
 - Naïve approach: check all edges, $O(mn)$ time in total
 - Better approach: keep track of distance of each node outside the growing component (S) from within S
 - Update when new vertex is added to S : only check edges going out from this new vertex



Prim's Algorithm

```
Prim-MST(G, s) {  
    /* starting vertex s  
       assume d(,) stores edge weights */  
    for each vertex v  
        D[v] := d(s, v) // initialise  
  
    S := {s}  
    while (S != V) {  
        find u in V - S with minimum D[u]  
        add u to S  
        for each v in V - S  
            D[v] := min(D[v], d(u, v))  
        }  
    }
```

Example Operation of Prim's

- Suppose start at f:


v	a	b	c	d	e	g
D[v]	11	∞	∞	∞	12	13

- Min. edge to a, add, update

v	a	b	c	d	e	g
D[v]	/	5	∞	∞	12	1

- Min. edge to g, add, update

v	a	b	c	d	e	g
D[v]	/	3	4	7	9	/



Prim's Example Operation (cont'd)


- Min. edge to b, add, update:

v	a	b	c	d	e	g
D[v]	/	/	4	7	9	/

- Min. edge to c, add, update:

v	a	b	c	d	e	g
D[v]	/	/	/	2	9	/

- ... and so on...







Running Time of Prim's

- Need two important operations:
 - *Find minimum*: $O(n)$ times
 - *Change value*: $O(m)$ times. Only consider the outgoing edges when a new vertex is added. $O(m)$ edges, and edges would not be considered again
- Using an array
 - $O(n)$ for finding minimum
 - $O(1)$ for change value
 - Total $O(m + n^2) = O(n^2)$ time
- Using a *heap*
 - $O(\log n)$ time for both operations
 - Total $O(m \log n + n \log n) = O(m \log n)$ time


Heaps

- Store elements in a *complete binary tree*
 - Left to right, do not start a new level unless it is full
 - *Balanced* (height = $O(\log n)$)
 - Parent \leq both children (for min-heaps; max-heap is opposite)
- Can easily be represented as an array $A[1..n]$, with all these in $O(1)$ time:
 - Parent($A[i]$) = $A[\lfloor i/2 \rfloor]$
 - Left-child($A[i]$) = $A[2i]$
 - Right-child($A[i]$) = $A[2i+1]$




Heap Operations

- Each in $O(\log n)$ time:
- (1) Decrease value
 - "Bubble up" (compare with parent and swap)
- (2) Insertion
 - Place at last available position, then similarly bubble up



Heap Operations

- (3) Delete minimum
 - Always at root
 - Remove root and replace it with last element in heap
 - "Sift down" the heap





Part 2: Shortest Paths




The Shortest Path Problem

- Many problems have similar nature
 - What is the shortest way to go from Leicester to London?
 - What is the quickest way to send a packet between two computers?
 - ...
- Abstract model
 - Given: a directed graph G with edge weights (e.g. distance), and a starting vertex s
 - Goal: find the shortest paths from s to all other vertices (*Single source shortest path*)
 - Path length = sum of edge weights
 - Not more difficult than just one s - t pair; same algorithms


Naïve Approaches

- Try #1: try all possible paths, calculate distance and find the minimum one
 - Problem: there can be exponentially many paths! (2^n)




- Try #2: just pick the shortest edge at each step (greedy!)

- Not correct



Optimal Substructure

- Consider the shortest path from s to t
 - Suppose it goes through x
 - Then this path from s to x must be a shortest path from s to x , too
 - (otherwise, we can replace it with a shorter path, the whole s - t path would also be shorter)




Also shortest path from s to x

Optimal Substructure


- Example:

- The path 1-7-1-5 is not optimal, because 1-7 can be replaced by 2-4




- Note that longest (simple) path does not have optimal substructure:

- a-b-c-d is a longest path from a to d
- but c-d is not longest path from c to d



A Similar Idea to BFS


- Recall the Breadth First Search algorithm:



- In fact, BFS is a shortest path algorithm if all edges have weight 1
 - Extend wavefront by 1 unit each time until reach target

Extending BFS...

- *Dijkstra's algorithm*: using a greedy idea
 - At every step, rather than extending the wavefront by 1 unit, we extend by the shortest "outgoing" edge
 - Example:





The Algorithm in Pseudocode

```
Dijkstra(G, s)  // source vertex s
  for each vertex v {
    D[v] := d(s, v) // provisional dist.
    Pred[v] := s if v is neighbor of s,
                  otherwise nil
  }
  S := {s}
  while (S != V) {
    find u in V - S with minimum D[u]
    add u to S
    for each v in V - S {
      if (D[u] + d(u,v) < D[v]) {
        D[v] := D[u] + d(u, v)
        Pred[v] := u
      }
    }
  }
}
```

Example Operation of Dijkstra's

- Shortest path from a to all nodes


- Initial

v	b	c	d	e	f	g
D[v]	5	∞	∞	∞	11	1
P[v]	a	/	/	/	a	a

- Shortest is to g

- Consider b: $D[g] + d(gb) < D[b]$
- Consider f: $D[g] + d(gf) > D[f]$

v	b	c	d	e	f	g
D[v]	4	5	8	10	11	1
P[v]	g	g	g	g	a	a



Dijkstra's Example (cont'd)


- Shortest is to b

v	b	c	d	e	f	g
D[v]	4	5	8	10	11	1
P[v]	g	g	g	g	a	a

- Shortest is to c

v	b	c	d	e	f	g
D[v]	4	5	7	10	11	1
P[v]	g	g	c	g	a	a

- ... and so on...



Getting the Actual Shortest Paths


- At the end,

v	b	c	d	e	f	g
D[v]	4	5	7	10	11	1
P[v]	g	g	c	g	a	a

← Shortest distance

← predecessor

- Can retrieve shortest path from this Pred[] array
- Example:



Shortest path tree



Running Time of Dijkstra's

- Note the similarity with Prim's algorithm for MST
 - Only difference is the update formula for $D[v]$
- Therefore, running time identical:
 - $O(n^2)$ for array
 - $O(m \log n)$ for heap
- Note that Dijkstra's algorithm only works for the case where *all edge weights are positive*